

Constraint-Graph Policy Optimizer for Constrained Multi-UAV Multi-Objective Path Planning

Anonymous Author^a

^aAnonymous for review, Review copy, Submission manuscript

ARTICLE INFO

Keywords:

multi-UAV path planning
multi-objective optimization
constrained evolutionary
optimization
constraint handling
benchmark protocol
hypervolume
feasibility diagnostics

ABSTRACT

Multi-UAV path planning is a constrained multi-objective optimization problem in which path quality, terrain clearance, kinematic feasibility, and inter-UAV separation must be considered together. Recent UAV planning and benchmarking studies emphasize realistic environment classes, fixed protocols, named baselines, statistical analysis, and reproducibility artifacts. However, many evolutionary UAV planners still treat constraint ranking, parent selection, and fleet-aware variation as weakly connected rules, making it difficult to inspect how observed violations influence search behavior. This paper presents the Constraint-Graph Policy Optimizer (CGPO), a lean graph-guided evolutionary optimizer for constrained multi-UAV path planning. CGPO builds a dynamic constraint-interaction graph over candidate fleet paths and uses the resulting tension signals in three coupled mechanisms: constraint-interaction graph construction (CIG), Pareto pressure-field selection (PPF), and orchestrated fleet variation (OVO). The active repository implementation contains no explicit repair stack, no graph-guided feasibility projection module, no constructive route-prior baseline, and no matched CGPO hybrid algorithms; all candidate improvement is performed through population-based evolutionary selection and variation with domain-bound clipping. A reduced matched evaluation on city, mountain, and suburban fleet-three cases uses population 40, 500 generations, five independent runs, matched safety constraints, hypervolume, IGD+, feasibility, mission diagnostics, and runtime. In this reduced protocol, CGPO is feasible in all 15 matched runs, achieves the best scenario-level HV rank, and outperforms NSGA-II, MOEA/D, and NMOPSO on matched run-level HV comparisons under Holm-corrected Wilcoxon tests. Focused inspection and ablation runs further verify fleet-size behavior and traceable CIG, PPF, and OVO switches. The contribution is a reproducible evolutionary graph-based optimizer and a journal-style reduced evaluation package for testing whether graph-derived constraint pressure improves constrained multi-UAV search.

1. Introduction

Unmanned aerial vehicles (UAVs) are increasingly used in tasks such as inspection, delivery, mapping, surveillance, and reconnaissance. Civil-application and path-planning surveys show that these missions require more than shortest geometric routes [24, 1]. Environment-complexity and evolutionary-computation surveys emphasize terrain, clearance, altitude, smoothness, and operational constraints [15, 14]. Recent UAV reviews and benchmark studies tie the same concerns to reproducible multi-objective evaluation [20, 6]. The difficulty increases further when several UAVs fly in a shared environment. Collaborative trajectory-planning and swarm-planning reviews emphasize coordination and separation constraints for multi-UAV missions [31, 16]. Multi-UAV optimization studies then turn those interactions into explicit constraint-handling and multitasking objectives [37, 19].

Multi-objective formulations are a natural fit for this setting because improving one criterion can degrade another. A path that shortens distance may pass closer to terrain or other UAVs. A path that increases clearance may require sharper turns or more energy.

Early comparative work already argued that UAV planners should be judged by problem-specific multi-objective indicators rather than one scalar route cost [3]. Decomposition-based and dimension-exploration UAV studies make the same trade-off structure explicit in constrained evolutionary settings [25, 33]. Urban many-objective planning and navigation-variable PSO studies further show that path quality, safety, and kinematics must be evaluated together [12, 9]. Recent constrained-UAV reviews and benchmark papers reinforce the need for safety diagnostics alongside objective scores [13, 6]. Multi-UAV and domain-specific operator studies add the fleet and operator-design pressures that motivate this work [19, 39]. The local domain-operator reference also motivates comparing CGPO against UAV-specific evolutionary mechanics rather than only generic MOEAs [5].

Despite this progress, constrained multi-UAV planning still has a control problem inside the optimizer. Evolutionary and swarm planners often combine archives, constraint ranking, crossover, and mutation through separate rules. This can produce useful results, but it weakens the connection between the constraints observed during search and the operators used to modify candidate paths. In

particular, a selection mechanism may suppress infeasible candidates without distinguishing harmful violations from useful near-boundary candidates, and a variation operator may alter one UAV path without considering its fleet-level coupling.

This paper addresses that gap with CGPO, a Constraint-Graph Policy Optimizer for constrained multi-UAV multi-objective path planning. CGPO uses one dynamic constraint-interaction graph to summarize terrain, altitude, turn-angle, pairwise-separation, and path-smoothing tension. The graph then drives parent-selection pressure and coordinated offspring generation. The current method is intentionally kept as a lean evolutionary algorithm: it does not use explicit repair, graph-guided feasibility projection, constructive A* route priors, or matched hybrid baselines. The intent is to make constraint pressure visible both to the optimizer and to the result analysis rather than to hide feasibility behind a post-processing step.

The key contributions are summarized as follows:

1. A constraint-interaction graph representation for fleet candidate paths that exposes terrain, altitude, turn-angle, pairwise-separation, and path-smoothing tension as recorded graph summaries.
2. A lean graph-guided evolutionary optimizer that couples constraint-interaction graph construction, Pareto pressure-field selection, and orchestrated fleet variation in one population-based search loop.
3. A reproducible SWEC-style evaluation protocol for constrained multi-UAV path planning, with fixed terrain cases, fleet sizes, budgets, baselines, safety diagnostics, hypervolume, IGD+, runtime, statistical analysis, and ablation variants.
4. A completed reduced matched comparison showing CGPO feasibility and HV advantages against NSGA-II, MOEA/D, and NMOPSO under the 40-population, 500-generation, five-run protocol, plus focused fleet-size and ablation inspections.

The remainder of the paper is organized into six sections. Section 2 reviews related work and formulates the benchmark problem. Section 3 describes CGPO. Section 4 gives the experimental protocol. Section 5 reports the reduced matched comparison, focused validation evidence, ablation results, and trace analysis. Section 7 summarizes the current evidence, reproducibility package, limitations, and data availability.

2. Related work and problem formulation

2.1. Related work

UAV path planning has been studied with graph search, sampling methods, potential fields, swarm intelligence, and evolutionary computation. Broad UAV surveys organize these techniques around mission requirements, environment complexity, and operational

safety [24, 1]. More recent surveys focus on environment complexity and evolutionary computation as recurring sources of route-planning difficulty [15, 14]. A recent UAV-planning review extends the same discussion to current algorithmic challenges and future directions [20]. For complex environments, multi-objective approaches are useful because decision makers often need a set of trade-off paths rather than one weighted-sum solution. Recent search-based multi-objective path-planning frameworks formalize this trade-off view beyond UAV-only settings [26]. Early comparative work emphasized the need for problem-specific quality indicators and statistical characterization when comparing multi-objective UAV planners [3]. More recent benchmark work frames UAV path planning around conflicting safety, efficiency, and regulatory constraints and introduces standardized urban, suburban, and mountainous scenarios [6]. A constrained-UAV mission-planning review reaches a similar conclusion from the perspective of constrained MOEA design [13].

Recent multi-UAV work strengthens the same direction. Two-fold constraint handling has been studied for multiple-UAV path planning [37]. Knowledge-assisted coevolution has been used for bi-objective multi-UAV planning [39]. Evolutionary multitasking has also been used for multiobjective multi-UAV path planning with adaptive operator selection and knowledge fusion [19]. Other local-corpus studies emphasize dimension-discrepancy search and urban many-objective planning for UAV routes [33, 12]. Navigation-variable PSO and domain-specific evolutionary operators provide two further UAV-specialized design points [9, 5]. These works motivate the present paper's fleet setting: single-UAV path quality is insufficient when inter-UAV constraints become part of the mission.

Benchmark papers in this area emphasize shared environments, fixed protocols, and clear baseline roles. Darlan et al. [6] explicitly argue that inconsistent scenarios and inaccessible code make algorithm comparison difficult. This repository follows that benchmark direction by using named protocol files, benchmark manifests, per-run artifacts, metrics summaries, and a separation between benchmark-safe algorithms and experimental research methods.

For CGPO, this matters because the algorithm is experimental. The manuscript therefore separates the method proposal from the comparative claim. The method can be described from code, but performance claims require the locked protocol outputs. The main, ablation, and single-UAV reference protocols are named YAML configs under `configs/`.

Constrained multi-objective optimization requires algorithms to handle objective trade-offs and constraint satisfaction simultaneously. General MOEA and constraint-handling surveys establish the core tension between convergence, diversity, and feasibility

Table 1

Online algorithm scope for the paper. Online denotes active paper scope, not dynamic real-time replanning.

Scope	Algorithm	Core mechanism	Paper role
Generic MOEA	NSGA-II	Pareto sorting and crowding distance	Standard MOEA baseline
Generic MOEA	MOEA/D	Weight-vector decomposition	Convergence-oriented baseline
UAV-specific MOEA	NMOPSO	Navigation-variable PSO with kinematic constraints	UAV path-planning baseline
UAV-specific MOEA	MOEA-2DE	Dimension exploration and discrepancy evolution	Single-UAV reference-code comparator
UAV-specific MOEA	EMMOP	Evolutionary multitasking and knowledge transfer	Single-UAV reference-code comparator
Constraint handling	NSGA-II with Deb rules	Feasibility-first constraint domination	Constrained-search baseline
Multi-UAV optimization	CCEA-ADVS	Co-evolutionary decomposition	Deferred fair-wiring target
Multi-UAV optimization	TSKAC-NSGA-II	Two-stage cooperative optimization	Runnable multi-UAV baseline
Separate MAPF	DTAPP-IICR	Delivery-time prioritized planning with iterative conflict resolution	Preflight coordination reference, reported separately

[38, 21]. Push-pull search and constrained-MOEA test-suite work show how feasibility pressure can be staged or benchmarked systematically [10, 18]. A more recent constrained-MOEA survey summarizes how archives, ranking rules, and auxiliary designs affect search behavior [17]. Multitasking-based constrained optimization studies show that knowledge transfer and constraint subsets can also influence constrained search [22, 35]. The CGPO design is closest to constraint-handling methods that treat infeasible and boundary candidates as informative rather than simply invalid.

The difference is representational. CGPO does not only compute a scalar constraint violation. It builds a constraint-interaction graph over route waypoints and uses the graph's tension field to guide selection pressure and variation. This lets the method record internal signals such as mean tension, pairwise edge count, smoothing edge count, conflict-cluster count, pressure entropy, perturbation scale, and coordinated clusters.

The paper uses an online/offline baseline scope. Here, online means that the method is part of the active comparison narrative or benchmark protocol; it does not mean real-time online replanning. Offline methods are retained as background, historical artifacts, or deferred comparators and are not used for headline numerical claims.

Table 1 lists the online algorithm set. NSGA-II is included as a canonical elitist multi-objective evolutionary algorithm [8]. MOEA/D is used as a decomposition-style convergence baseline. NMOPSO represents the UAV path-planning PSO family, grounded in MOPSO and SMPSO foundations [4, 23] and recent navigation-variable UAV PSO work [9]. MOEA-2DE and EMMOP represent UAV-specialized external reference-code methods. TSKAC-NSGA-II represents a UAV-specialized multi-UAV coevolutionary baseline related to knowledge-assisted planning [39]. CCEA-ADVS is kept in the MOEA-style online paper scope as a multi-UAV coevolution reference. DTAPP-IICR is kept as a separate MAPF/preflight coordination baseline because it returns a coordinated fleet plan rather than a Pareto population.

All other algorithms are offline for the current manuscript scope. This includes L-SHADE-CDP, C-TAEA, GCNMOEA, MO-MFEA-II, MOEA-DWA-ASTAR, the legacy SAC/relational-state variants, and custom MOGWO variants. They may be discussed as related work or future comparator targets, but no headline numeric comparison is reported for them in this manuscript unless the protocol is explicitly expanded.

2.2. Problem formulation

The benchmark represents an environment as a terrain map with horizontal bounds, altitude limits, and scenario-specific obstacle or terrain structure. The active CGPO/SWEC protocol uses the paper_medium scenario set with three named cases: c_100, m_100, and s_120. These correspond to the city, mountain, and suburban terrain families used by the repository's benchmark workflow.

A fleet candidate is a set of waypoint paths

$$X = \{P_1, \dots, P_m\}, \quad P_u = (x_{u,i}, y_{u,i}, z_{u,i})_{i=1}^n, \quad (1)$$

where m is the fleet size and each path has fixed start and goal endpoints. CGPO modifies internal waypoints only. This keeps the optimization focused on route shape and fleet interaction rather than changing the mission definition.

The repository evaluator reports four objective channels through the common UAV benchmark pipeline. The manuscript intentionally does not rename these channels into unsupported physical quantities beyond the repository documentation; they are treated as path-quality, terrain/altitude, smoothness or mission-quality, and safety-related objective channels. In addition, the metrics pipeline reports feasible ratio, run success, hypervolume (HV), IGD+, runtime, minimum clearance, maximum turn angle, mission conflict rate, mission makespan, and mission energy.

HV and IGD+ measure Pareto-set quality, while the mission diagnostics prevent unsafe paths from being hidden behind aggregate objective scores. The use of multiple indicators follows established multiobjective performance-assessment guidance [40]. Hypervolume is

Table 2

Mission constraints fixed by the registered protocol.

Parameter	Value	Evidence source
separationMin	10.0	Head-to-head config
maxTurnDeg	75.0	Head-to-head config
safeDist	20.0	Head-to-head config
droneSize	1.0	Head-to-head config
hardCollisionConstraint	true	Head-to-head config

included because it is a standard Pareto-set quality indicator with dedicated algorithmic treatment [32]. UAV planner comparison work also supports reporting problem-specific diagnostics rather than relying on one aggregate indicator [3]. This follows the evidence style of SWEC UAV papers, where tables, convergence, path visualizations, and safety diagnostics are used together rather than relying on one metric.

The registered protocol fixes the mission constraints shown in Table 2. Feasibility is determined by the shared mission evaluator and constraint-violation machinery used by the benchmark. The main hard constraints are terrain clearance, altitude bounds, turn-angle limits, and inter-UAV separation. The same evaluator is used for CGPO and non-CGPO baselines in the benchmark workflow.

3. Method: Constraint-Graph Policy Optimizer

3.1. Overview

The central design decision in CGPO is to make constraint pressure inspectable before it is used. Terrain, altitude, turn, pairwise, and path-smoothing pressures are converted into a graph-based tension field. The optimizer then uses this field in two places: selection and variation. CGPO deliberately avoids explicit repair and graph-guided feasibility projection in the active method, so any observed improvement must come from evolutionary search pressure rather than from a separate correction layer. If a mechanism does not help, the ablation protocol should expose it.

3.2. Graph-guided evolutionary mechanisms

The constraint-interaction graph (CIG) is implemented in `uav_benchmark/algorithms/cgpo/cig.py`. For each candidate fleet path, valid internal waypoints become graph nodes with three-dimensional tension vectors. The implementation adds terrain and altitude pressure when relative altitude approaches or violates the safe clearance envelope, turn pressure when horizontal turn geometry approaches or exceeds the configured turn limit, pairwise pressure when UAV waypoints are close within a temporal window, and a small path-smoothing pressure when edge coupling is enabled. The smoothing-edge weight is a fixed scalar

and is decoupled from the PPF objective weights, so the same graph construction is used regardless of the selection state.

The graph returns the full tension tensor and summary values: scalar tension, per-UAV tension, mean tension, maximum tension, terrain-edge count, turn-edge count, pairwise-edge count, smoothing-edge count, conflict-cluster count, and active constraint types. These trace values are not final evidence by themselves. They are used to interpret external results and ablations.

The Pareto pressure field (PPF) is implemented in `uav_benchmark/algorithms/cgpo/ppf.py`. PPF combines non-dominated sorting rank, constraint violation, crowding-distance diversity, boundary usefulness, and graph tension. When enabled, it produces parent-selection probabilities, feasibility pressure, exploration scale, objective weights, boundary mass, feasible mass, pressure entropy, and stratum counts.

This design reflects a common constrained-optimization observation: not every infeasible candidate is equally useless. A near-boundary candidate with useful geometry may yield a high-quality feasible offspring under graph-aware variation, while a feasible but crowded candidate may contribute little diversity. PPF makes this distinction explicit in the selection layer rather than relying on a downstream correction stage to convert near-boundary candidates after the fact.

The orchestrated variation operator (OVO) is implemented in the CGPO module. OVO blends parent waypoints using inverse graph tension, perturbs internal waypoints according to the PPF exploration scale, and applies a waypoint-aligned fleet-coordination pass that shifts nearby UAV waypoints apart at the same temporal index when their pairwise separation is violated. The operator records perturbation scale, number of perturbed waypoints, coordinated cluster count, and parent-blend entropy.

The coordination pass is the main difference between a generic path mutation and a multi-UAV operator. It uses the same waypoint indexing as the pairwise CIG edges, so the coordination decision is consistent with the constraint signals that selection has already seen, and the operator avoids the higher-cost trajectory-resampling scan used in the earlier prototype.

3.3. Search loop and ablation hooks

The CGPO runner is implemented in `uav_benchmark/algorithms/cgpo/__init__.py`. It initializes a population with uniformly random fleet candidates inside the domain bounds (no constructive A*-corridor seeds and no analytic ramp seeds, so the published method has no terrain-aware initialization shortcut), evaluates candidates with the shared mission evaluator, builds CIG summaries, computes PPF, samples parent




Figure 1: CGPO search loop and active mechanism coupling. CIG converts constraint pressures into graph tension summaries, PPF turns Pareto and constraint information into parent-selection pressure, and OVO uses those signals for fleet-aware offspring generation.

Table 3

CGPO three-mechanism ablation variants supported by the active repository implementation.

Variant	Removed mechanism
CGPO_full	none
CGPO_random_only	all three (CIG, PPF, OVO)
CGPO_no_cig_edge_coupling	CIG edge coupling
CGPO_no_ppf_pressure	PPF pressure
CGPO_no_ovo_variation	OVO variation operator
CGPO_no_ovo_fleet_coordination	OVO fleet-coordination pass only

pairs from PPF probabilities, applies OVO, projects offspring back into the domain bounds, evaluates offspring, and updates the constrained elite archive using NSGA-II constraint-domination semantics. Final artifacts are saved using the common benchmark result format.

The conceptual ablation variants are registered in `uav_benchmark/benchmark.py` and summarized in Table 3. A mechanism is considered supported only if the ablation shows a meaningful external degradation after removal or a trace-backed explanation for the observed behavior.

4. Experimental protocol

The reduced matched comparison uses CGPO, NSGA-II, MOEA/D, and NMOPSO as MOEA-style optimization rows. The run budget is five independent runs, 500 generations, requested population 40, seed 11, fleet mode, fleet size 3, the `c_100`, `m_100`, and `s_120` problems, the `paper_medium` scenario set, GPU mode off, metric interval 10, and `nRep`: 40. MOEA/D's NBI weight generation uses 35 internal weight vectors for the requested 40-population budget; this is recorded as a protocol note rather than hidden. Larger-scope methods such as MOEA-2DE, EMMOP, CCEA-ADVS,

TSKAC-NSGA-II, and DTAPP-IICR remain part of the planned full protocol, but they are not used for reduced-budget claims in this manuscript. The protocol is included because multiobjective comparison requires transparent indicator reporting [40]. It also follows recent UAV benchmark practice by fixing scenarios, budgets, baselines, and problem definitions [6]. The multi-UAV and domain-operator references motivate keeping fleet-aware and UAV-specialized baselines in the broader comparison set [19, 5].

The separate online reference-code comparison is defined by the single-UAV reference protocol. It includes MOEA-2DE and EMMOP with CGPO on fleet size 1. These methods are not mixed into the main fleet table because they preserve author-code mechanics or single-UAV assumptions. CCEA-ADVS remains an online paper-scope reference but offline in the runnable benchmark until a fair local implementation is available.

The full ablation is defined by the CGPO ablation protocol. It uses the same terrain cases, fleet sizes, runs, generations, population, constraints, and metrics as the main comparison, but replaces the algorithm list with the five CGPO variants in Table 3.

The comparative analysis should report mean, standard deviation, median, and interquartile range where available. HV and IGD+ are the primary Pareto-set quality metrics. Feasible ratio, run success, mission conflict rate, minimum clearance, maximum turn angle, makespan, energy, and runtime are mandatory diagnostics.

For statistical testing, the planned analysis is matched non-parametric comparison where pairing is valid, with Holm correction for multiple comparisons. Statistical labels are used only when the corresponding pairwise or ablation result artifacts support them.

Matched reduced CGPO/SWEC comparison: 40 population, 500 generations, 5 runs, fleet size 3




Figure 2: Data-derived reduced comparison across the three fleet-three scenarios. Bars aggregate scenario-level means from results/cgpo_swec_matched_40x500_5run/metrics/benchmark_metrics_summary.csv.

Table 4

Reduced matched comparison across c_100_uav3, m_100_uav3, and s_120_uav3. HV and runtime are averaged over scenario-level means; $n = 5$ runs per scenario.

Role	Method	Runs	FR	HV	IGD+	Conflict	Runtime
Proposed	CGPO	15	1.0000	0.1728	0.0981	0.0000	76.4872
Decomposition baseline	MOEA/D	15	0.6000	0.0834	0.2382	0.0000	24.0834
Canonical MOEA baseline	NSGA-II	15	0.0000	0.0000	n/a	0.0000	25.7643
UAV PSO baseline	NMOPSO	15	0.0000	0.0000	n/a	0.0001	26.2931

5. Results and ablation

5.1. Reduced matched comparison

The publication-minimum comparison in Table 4 uses configs/cgpo_swec_matched_40x500_5run.yaml. It compares CGPO with NSGA-II, MOEA/D, and NMOPSO on the city, mountain, and suburban fleet-three cases. Each cell aggregates five runs per scenario, so the table should be read as reduced-budget evidence rather than as the complete 30-run protocol.

CGPO is the only method that returns feasible populations in every matched run. It has the best scenario-level HV rank on all three fleet-three scenarios and wins every available matched run-level HV comparison against the three reduced baselines. MOEA/D is the strongest non-CGPO baseline in this reduced protocol, but it fails the city fleet-three case and uses 35 NBI weight vectors for the requested population 40. NSGA-II and NMOPSO do not produce feasible final populations under the same budget, which makes their HV zero after feasible-mask filtering. These findings support a reduced-budget claim that CGPO is more reliable under the tested constrained fleet setting; they do not yet replace the larger planned protocol with UAV-specialized and MAPF baselines.

5.2. Validation runs

Two single-run full-budget inspections are available from the focused table protocol configs/cgpo_swec_inspection_40x500_1run.yaml. Both use 500 generations, population 40, seed 11, GPU mode off, s_120, hard collision constraints, separationMin =

10.0, maxTurnDeg = 75.0, safeDist = 20.0, droneSize = 1.0, and nRep = 40. They are useful for implementation and feasibility inspection, but they are not statistically meaningful because each has only one run and no baseline counterpart.

These runs show that the current CGPO implementation can complete full-budget s_120 fleet inspections and produce feasible reported outputs in the tested seeds. Under the 40-population table protocol, the three-UAV run achieved HV 0.3006 and the eight-UAV run achieved HV 0.0442; both reported feasible ratio 1.0, run success 1.0, and zero mission conflict rate.

The runtime difference between the two inspections is consistent with the fleet-size burden: the s_120_uav8 inspection took 214.1621 seconds compared with 74.2049 seconds for s_120_uav3. This observation is an implementation diagnostic rather than a formal complexity result, because it is based on one run per fleet size.

5.3. Ablation and trace analysis

Table 7 reports a focused one-run ablation on s_120_uav3. Each variant uses 500 generations, population 40, seed 11, GPU mode off, and the same mission constraints as the validation table. The results are real benchmark outputs, but they remain descriptive because there is only one run per variant.

The trace harness records, per generation and per variant, mean and per-UAV CIG tension, terrain-edge count, turn-edge count, pairwise-edge count, smoothing-edge count, conflict-cluster count, PPF

Table 5

Matched run-level HV comparison against CGPO. Wilcoxon tests use the alternative hypothesis that CGPO has larger HV; p_{Holm} applies Holm correction across the three baseline comparisons.

Comparison	Matched n	W/T/L	Mean ΔHV	Median ΔHV	p_{Holm}
CGPO vs. MOEA/D	14	14/0/0	0.0841	0.0767	9.16×10^{-5}
CGPO vs. NSGA-II	15	15/0/0	0.1731	0.1704	9.16×10^{-5}
CGPO vs. NMOPSO	15	15/0/0	0.1731	0.1704	9.16×10^{-5}

Table 6

Completed full-budget CGPO validation runs. These values support implementation and feasibility inspection; broader comparative claims require the complete multi-run protocol.

Directory	Problem	Runs	FR	Success	HV	IGD+	Runtime	Clearance	Turn	Conflict	Energy
cgpo_swec_inspection_40x500_1run	s_120_uav3	1	1.0	1.0	0.3006	0.0000	74.2049	3.7876	74.9551	0.0	0.1192
cgpo_swec_inspection_40x500_1run	s_120_uav8	1	1.0	1.0	0.0442	0.0000	214.1621	1.3278	61.3503	0.0	0.0701

feasibility pressure, exploration scale, boundary mass, feasible mass, pressure entropy, OVO perturbation scale, perturbed waypoint count, coordinated cluster count, and parent-blend entropy. The expected, code-pinned response of these signals to the ablation switches is: when CIG edge coupling is disabled, pairwise and smoothing edge counts drop to zero and PPF tension input flattens; when PPF pressure is disabled, parent selection becomes uniform and the feasibility-pressure and pressure-entropy series degenerate; when OVO variation is disabled, the perturbation scale and parent-blend entropy series drop to zero; and when only the OVO fleet-coordination pass is disabled, the coordinated-cluster count drops to zero while the rest of the OVO trace remains active. The full ablation analysis script verifies this code-pinned mapping and reports it alongside the external metrics so that any external degradation can be linked back to the correct internal mechanism.

Ablations should be interpreted by external metrics first and traces second. In the reduced one-run ablation, disabling OVO variation clearly reduces HV and worsens IGD+, while disabling CIG edge coupling or PPF pressure does not degrade this single s_120_uav3 case. These results make OVO variation the best-supported component in the reduced evidence and make CIG/PPF conditional mechanisms that need broader multi-scenario ablation before stronger necessity claims. If a trace changes but the external metrics do not improve, the mechanism should be presented as inspectable but not necessary. If a removed component

improves external metrics, the method should be simplified rather than defended by terminology.

Table 7

Focused 40-population, 500-generation, one-run ablation on s_120_uav3. Values are descriptive single-run results.

Variant	Runs	FR	HV mean	IGD+ mean	Runtime	Clearance	Turn	Conflict
CGPO_full	1	1.0	0.4540	0.0514	73.7439	3.7876	74.9551	0.0
CGPO_random_only	1	1.0	0.3676	0.1927	47.3363	3.8462	66.2798	0.0
CGPO_no_cig_edge_coupling	1	1.0	0.4590	0.0480	60.6004	1.4211	71.0968	0.0
CGPO_no_ppf_pressure	1	1.0	0.4639	0.0358	74.3816	3.7876	46.9165	0.0
CGPO_no_ovo_variation	1	1.0	0.3182	0.2129	69.0420	3.5678	74.5429	0.0
CGPO_no_ovo_fleet_coordination	1	1.0	0.4479	0.0446	71.5178	3.8462	59.0840	0.0

6. Reproducibility and limitations

The repository contains the configs and result-artifact conventions needed for reproduction. The reduced matched comparison, focused fleet-size inspection, and focused ablation are controlled by the three CGPO/SWEC 40x500 YAML configs under `configs/`. The matched-comparison command writes benchmark manifests, raw per-run artifacts, metric summaries, pairwise statistics, rank summaries, table CSVs, and the data-derived result figure.

The result bundle should include benchmark manifests, config files, raw per-run artifacts, metric summaries, pairwise statistics, win/tie/loss summaries, ablation summaries, trace summaries, figure source data, and the final manuscript.

The completed 40-population inspection manifest records macOS arm64, Python 3.10.19, NumPy 2.2.6, SciPy 1.15.3, repository branch `codex/repo-hygiene-and-fixes`, commit `8d30d19d`, and a dirty working tree. A clean submission should record the final commit hash and whether the result directory was generated from a clean or dirty tree.

The study has four important limitations. First, the completed numerical evidence is a reduced comparative benchmark, not the full planned protocol. It includes a five-run matched comparison against NSGA-II, MOEA/D, and NMOPSO on three fleet-three scenarios, plus focused single-run fleet-size and ablation inspections. This supports reduced-budget claims against these baselines, but it does not support broad state-of-the-art claims across UAV-specialized or MAPF-style methods. MOEA/D also generated 35 NBI weight vectors for the requested population of 40, so that detail should be retained when interpreting its row.

Second, the benchmark is simulation-centered and static. It does not model online replanning, wind, communication loss, onboard controller dynamics, hardware-in-the-loop validation, or real-flight execution.

Third, equal generations and population do not guarantee perfectly equal objective-evaluation counts for all algorithms. The lean three-mechanism CGPO uses one full mission evaluation per offspring and no projection-proxy evaluations. The paper records candidate evaluations and total mission-evaluation calls in run metadata and reports these counts next to runtime whenever CGPO is compared with baselines.

Fourth, CGPO has computational overhead. The 40-population full-budget single-run inspection on `s_120_uav8` took 214.1621 seconds on the recorded local environment. Runtime must therefore remain part of the final quality discussion.

7. Conclusion

This paper presents CGPO, a constraint-graph policy optimizer for constrained multi-UAV multi-objective path planning. The published method uses a dynamic constraint-interaction graph to connect Pareto pressure-field selection and orchestrated fleet-aware variation, with no explicit feasibility projection, repair stack, constructive route priors, or matched hybrid baselines in the active repository implementation. The repository exposes the relevant control paths and trace variables, and the SWEC-style protocol specifies the terrain cases, baselines, budgets, metrics, statistics, and ablations needed for a defensible evaluation.

The completed reduced-budget artifacts show that CGPO returns feasible populations in all matched fleet-three runs and outperforms NSGA-II, MOEA/D, and NMOPSO on matched HV comparisons under the tested budget. The focused fleet-size inspection confirms that CGPO can complete `s_120` three- and eight-UAV runs with zero reported mission conflict rate. The one-run ablation supports OVO variation as the clearest contributor in the tested case, while CIG and PPF remain conditional mechanisms requiring broader ablation evidence.

The manuscript is based on the local UAV benchmark repository. The active reduced-budget paper configs are the matched comparison, inspection, and ablation CGPO/SWEC YAML files under `configs/`. The CGPO implementation is under `uav_benchmark/algorithms/cgpo/`. The completed numerical artifacts are stored in the corresponding CGPO/SWEC result directories and summarized in the paper-artifact manifest.

During manuscript preparation, AI-assisted writing tools were used to help organize and edit text. All scientific claims, numerical values, and artifact references were checked against local repository files and result outputs by the authors.

References

- [1] Aggarwal, S., Kumar, N., 2020. Path planning techniques for unmanned aerial vehicles: A review, solutions, and challenges. *Computer Communications* 149, 270–299. URL: <https://doi.org/10.1016/j.comcom.2019.10.014>, doi:10.1016/j.comcom.2019.10.014.
- [2] Bali, K.K., Gupta, A., Ong, Y.S., Tan, P.S., 2021. Cognizant multitasking in multiobjective multifactorial evolution: MO-MFEA-II. *IEEE Transactions on Cybernetics* 51, 1784–1796. doi:10.1109/TCYB.2020.2981733.
- [3] Besada-Portas, E., de la Torre, L., Moreno, A., Risco-Martín, J.L., 2013. On the performance comparison of multi-objective evolutionary UAV path planners. *Information Sciences* 238, 111–125. doi:10.1016/j.ins.2013.02.022.
- [4] Coello Coello, C.A., Lechuga, M.S., 2002. MOPSO: A proposal for multiple objective particle swarm optimization, in: *Proc. IEEE Congress on Evolutionary Computation (CEC)*, pp. 1051–1056. URL: <https://ieeexplore.ieee.org/document/1004388>, doi:10.1109/CEC.2002.1004388.

- [5] Darlan, D., Ajani, O.S., Mallipeddi, R., 2026. Evolutionary algorithm with domain-specific operators for UAV path planning. *Swarm and Evolutionary Computation* 101, 102267. URL: <https://doi.org/10.1016/j.swevo.2025.102267>, doi:10.1016/j.swevo.2025.102267.
- [6] Darlan, D., Ajani, O.S., Paul, A., Mallipeddi, R., 2025. A multi-objective benchmark for UAV path planning with baseline results. *Swarm and Evolutionary Computation* 96, 101968. doi:10.1016/j.swevo.2025.101968.
- [7] Deb, K., Jain, H., 2014. An evolutionary many-objective optimization algorithm using reference-point-based nondominated sorting approach, Part I. *IEEE Transactions on Evolutionary Computation* 18, 577–601. URL: <https://ieeexplore.ieee.org/document/6600851>, doi:10.1109/TEVC.2013.2281535.
- [8] Deb, K., Pratap, A., Agarwal, S., Meyarivan, T., 2002. A fast and elitist multiobjective genetic algorithm: NSGA-II. *IEEE Transactions on Evolutionary Computation* 6, 182–197. doi:10.1109/4235.996017.
- [9] Duong, T.T.N., Bui, D.N., Phung, M.D., 2025. Navigation variable-based multi-objective particle swarm optimization for UAV path planning with kinematic constraints. *Neural Computing and Applications* doi:10.1007/s00521-024-10945-1.
- [10] Fan, Z., Li, W., Cai, X., Li, H., Wei, C., Zhang, Q., Deb, K., Goodman, E., 2019. Push and pull search for solving constrained multi-objective optimization problems. *Swarm and Evolutionary Computation* 44, 665–679. URL: <https://doi.org/10.1016/j.swevo.2018.08.017>, doi:10.1016/j.swevo.2018.08.017.
- [11] Gupta, A., Ong, Y.S., Feng, L., 2016. Multifactorial evolution: Toward evolutionary multitasking. *IEEE Transactions on Evolutionary Computation* 20, 343–357. doi:10.1109/TEVC.2015.2458037.
- [12] Hohmann, M., et al., 2023. Three-dimensional urban path planning for aerial vehicles regarding many objectives. *IEEE Open Journal of Intelligent Transportation Systems* URL: <https://doi.org/10.1109/OJITS.2023.3299496>, doi:10.1109/OJITS.2023.3299496.
- [13] Huang, G., Hu, M., Yang, X., Wang, X., Wang, Y., Huang, F., 2024. A review of constrained multi-objective evolutionary algorithm-based unmanned aerial vehicle mission planning: Key techniques and challenges. *Drones* 8, 316. URL: <https://doi.org/10.3390/drones8070316>, doi:10.3390/drones8070316.
- [14] Jiang, L., et al., 2024. Evolutionary computation for unmanned aerial vehicle path planning: A survey. *Artificial Intelligence Review* 57, 109. doi:10.1007/s10462-024-10913-0.
- [15] Jones, M., et al., 2023. Path-planning for unmanned aerial vehicles with environment complexity considerations: A survey. *ACM Computing Surveys* 55. URL: <https://doi.org/10.1145/3570723>, doi:10.1145/3570723.
- [16] Li, J., Li, J., Zhang, J., Meng, W., 2026. A comprehensive review of path-planning algorithms for multi-UAV swarms. *Drones* 10, 11. URL: <https://doi.org/10.3390/drones10010011>, doi:10.3390/drones10010011.
- [17] Liang, J., Ban, X., Yu, K., Qu, B., Qiao, K., Yue, C., Chen, K., Tan, K.C., 2023. A survey on evolutionary constrained multiobjective optimization. *IEEE Transactions on Evolutionary Computation* 27, 201–221. URL: <https://ieeexplore.ieee.org/document/9723472>, doi:10.1109/TEVC.2022.3155533.
- [18] Ma, Z., Wang, Y., 2019. Evolutionary constrained multi-objective optimization: Test suite construction and performance comparisons. *IEEE Transactions on Evolutionary Computation* 23, 972–986. URL: <https://ieeexplore.ieee.org/document/8632683>, doi:10.1109/TEVC.2019.2896967.
- [19] Meng, K., Wu, B., Xin, B., Deng, F., Chen, C., 2025a. Multiobjective multi-UAV path planning via evolutionary multitasking optimization with adaptive operator selection and knowledge fusion. *Swarm and Evolutionary Computation* 99, 102145. doi:10.1016/j.swevo.2025.102145.
- [20] Meng, W., Zhang, X., Zhou, L., Guo, H., Hu, X., 2025b. Advances in UAV path planning: A comprehensive review of methods, challenges, and future directions. *Drones* 9, 376. URL: <https://doi.org/10.3390/drones9050376>, doi:10.3390/drones9050376.
- [21] Mezura-Montes, E., Coello Coello, C.A., 2011. Constraint-handling in nature-inspired numerical optimization: Past, present and future. *Swarm and Evolutionary Computation* 1, 173–194. doi:10.1016/j.swevo.2011.10.001.
- [22] Ming, F., Gong, W., Wang, L., Gao, L., 2024. Constrained multiobjective optimization via multitasking and knowledge transfer. *IEEE Transactions on Evolutionary Computation* 28, 77–88. doi:10.1109/TEVC.2022.3230822.
- [23] Nebro, A.J., Durillo, J.J., Garcia-Nieto, J., Coello Coello, C.A., Luna, F., Alba, E., 2009. SMP-PSO: A new PSO-based metaheuristic for multi-objective optimization, in: *Proc. IEEE Symposium on Computational Intelligence in Multi-Criteria Decision-Making*, pp. 66–73. URL: <https://ieeexplore.ieee.org/document/4938830>, doi:10.1109/MCDM.2009.4938830.
- [24] Otto, A., Agatz, N., Campbell, J., Golden, B., Pesch, E., 2018. Optimization approaches for civil applications of unmanned aerial vehicles (UAVs) or aerial drones: A survey. *Networks* 72, 411–458. URL: <https://doi.org/10.1002/net.21818>, doi:10.1002/net.21818.
- [25] Peng, C., Qiu, S., 2022. A decomposition-based constrained multi-objective evolutionary algorithm with a local infeasibility utilization mechanism for UAV path planning. *Applied Soft Computing* 118, 108495. doi:10.1016/j.asoc.2022.108495.
- [26] Ren, Z., et al., 2025. EMOA*: A framework for search-based multi-objective path planning. *Artificial Intelligence* 339, 104260. doi:10.1016/j.artint.2024.104260.
- [27] Storn, R., Price, K., 1997. Differential evolution – a simple and efficient heuristic for global optimization over continuous spaces. *Journal of Global Optimization* 11, 341–359. doi:10.1023/A:1008202821328.
- [28] Tanabe, R., Fukunaga, A., 2013. Success-history based parameter adaptation for differential evolution, in: *Proc. IEEE Congress on Evolutionary Computation (CEC)*, pp. 71–78. URL: <https://ieeexplore.ieee.org/document/6557555>, doi:10.1109/CEC.2013.6557555.
- [29] Tanabe, R., Fukunaga, A.S., 2014. Improving the search performance of SHADE using linear population size reduction, in: *Proc. IEEE Congress on Evolutionary Computation*, pp. 1658–1665. doi:10.1109/CEC.2014.6900380.
- [30] Tian, Y., Cheng, R., Zhang, X., Jin, Y., 2017. PlatEMO: A MATLAB platform for evolutionary multi-objective optimization. *IEEE Computational Intelligence Magazine* 12, 73–87. doi:10.1109/MCI.2017.2742868.
- [31] Wang, L., Huang, W., Li, H., Li, W., Chen, J., Wu, W., 2024. A review of collaborative trajectory planning for multiple unmanned aerial vehicles. *Processes* 12, 1272. URL: <https://doi.org/10.3390/pr12061272>, doi:10.3390/pr12061272.
- [32] While, L., Hingston, P., Barone, L., Huband, S., 2006. A faster algorithm for calculating hypervolume. *IEEE Transactions on Evolutionary Computation* 10, 29–38. URL: <https://ieeexplore.ieee.org/document/1583625>, doi:10.1109/TEVC.2005.851275.
- [33] Xu, S., et al., 2024. A multi-objective evolutionary algorithm based on dimension exploration and discrepancy evolution for UAV path planning problem. *Information Sciences* 657, 119977. URL: <https://doi.org/10.1016/j.ins.2023.119977>, doi:10.1016/j.ins.2023.119977.
- [34] Yan, P., Tian, Y., Liu, Y., 2025. An indicator-based multi-objective evolutionary algorithm assisted by improved

- graph convolutional networks. *Swarm and Evolutionary Computation* 94, 101892. URL: <https://doi.org/10.1016/j.swevo.2025.101892>, doi:10.1016/j.swevo.2025.101892.
- [35] Yu, K., Wang, L., Liang, J., Wang, H., Qiao, K., Liang, T., 2024. Constraint subsets-based evolutionary multitasking for constrained multiobjective optimization. *Swarm and Evolutionary Computation* 87, 101531. doi:10.1016/j.swevo.2024.101531.
- [36] Zhang, J., Sanderson, A.C., 2009. JADE: Adaptive differential evolution with optional external archive. *IEEE Transactions on Evolutionary Computation* 13, 945–958. URL: <https://ieeexplore.ieee.org/document/5208221>, doi:10.1109/TEVC.2009.2014613.
- [37] Zhang, W., et al., 2024. A novel multi-objective evolutionary algorithm with a two-fold constraint-handling mechanism for multiple UAV path planning. *Expert Systems with Applications* 238, 121862. doi:10.1016/j.eswa.2023.121862.
- [38] Zhou, A., Qu, B.Y., Li, H., Zhao, S.Z., Suganthan, P.N., Zhang, Q., 2011. Multiobjective evolutionary algorithms: A survey of the state of the art. *Swarm and Evolutionary Computation* 1, 32–49. URL: <https://doi.org/10.1016/j.swevo.2011.03.001>, doi:10.1016/j.swevo.2011.03.001.
- [39] Zhou, T., Zhou, Z., Qiu, H., Niu, B., Yue, G.X.G., Pedrycz, W., 2024. Two-stage knowledge-assisted coevolutionary NSGA-II for bi-objective path planning of multiple unmanned aerial vehicles. *Swarm and Evolutionary Computation* 90, 101680. doi:10.1016/j.swevo.2024.101680.
- [40] Zitzler, E., Thiele, L., Laumanns, M., Fonseca, C.M., Grunert da Fonseca, V., 2003. Performance assessment of multiobjective optimizers: An analysis and review. *IEEE Transactions on Evolutionary Computation* 7, 117–132. URL: <https://ieeexplore.ieee.org/document/1197687>, doi:10.1109/TEVC.2003.810758.